

# “Deployment for free”: removing the need to write model deployment code at Stitch Fix

Stanford CS329S February 2022

**Stefan Krawczyk**

 @stefkrawczyk  
 [linkedin.com/in/skrawczyk](https://www.linkedin.com/in/skrawczyk)

Try out Stitch Fix → [goo.gl/Q3tCQ3](https://goo.gl/Q3tCQ3)

#CS329S #MLOps #machinelearning

STITCH FIX



> Stitch Fix

“Deployment for free”

Model Envelope & envelope mechanics

Impact of being on-call

Summary & Future Work

# Stitch Fix is a personal styling service

Key points:

1. Very algorithmically driven company
2. Single DS Department: Algorithms (135+)
3. “Full Stack Data Science”
  - a. No reimplementation handoff
  - b. End to end ownership
  - c. Built on top of data platform tools & abstractions.

For more information: <https://algorithms-tour.stitchfix.com/> & <https://cultivating-algos.stitchfix.com/>

# Where do I fit in?



Pre-covid look

Stefan Krawczyk


Mgr. Data Platform - Model Lifecycle



MSCS'10



iDIBON



# STITCH FIX

Checkout out our open source dataflow library that helps manage feature/workflow code for you:

- <https://github.com/stitchfix/hamilton/>

A large, faint, circular watermark of the Stitch Fix logo is visible on the left side of the slide. It features a stylized 'S' and 'F' intertwined within a circle.

Stitch Fix

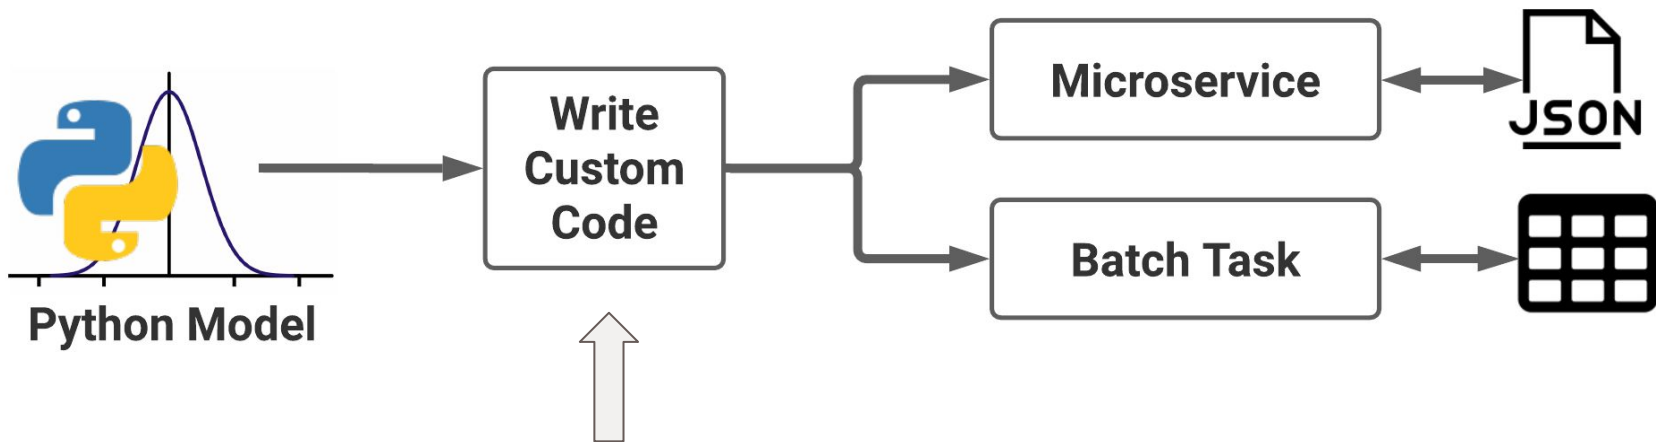
> “Deployment for free”

Model Envelope & envelope mechanics

Impact of being on-call

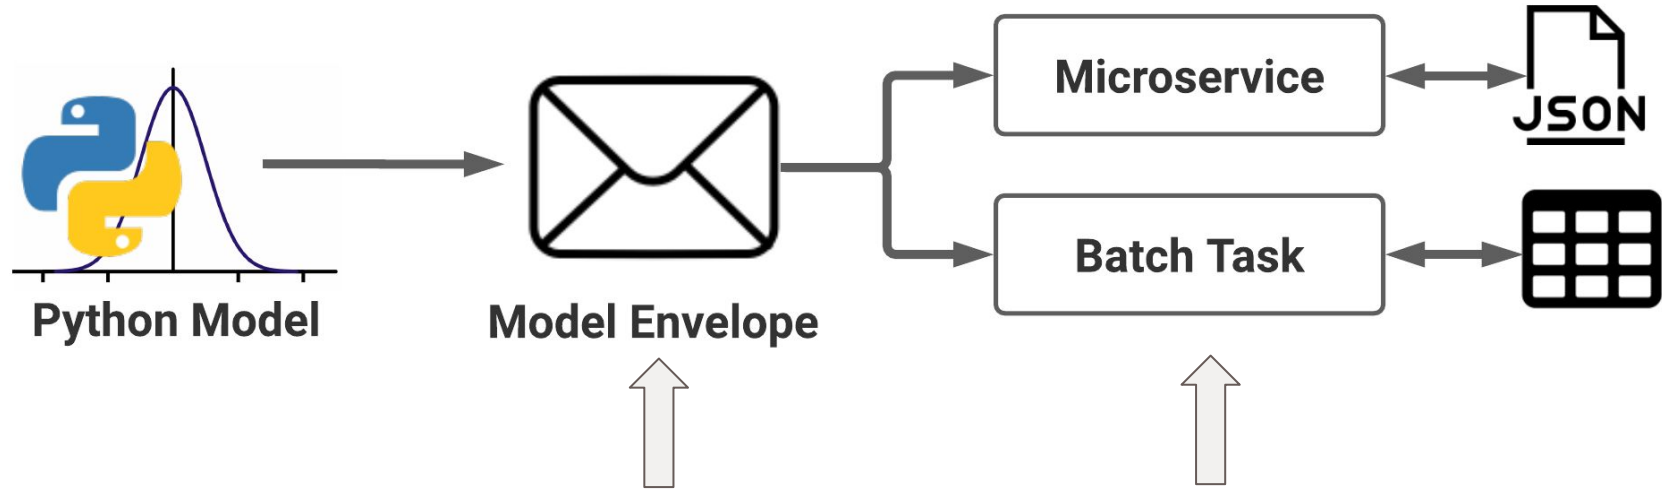
Summary & Future Work

# Typical Model Deployment Process



- Many ways to approach.
- Heavily impacts MLOps.

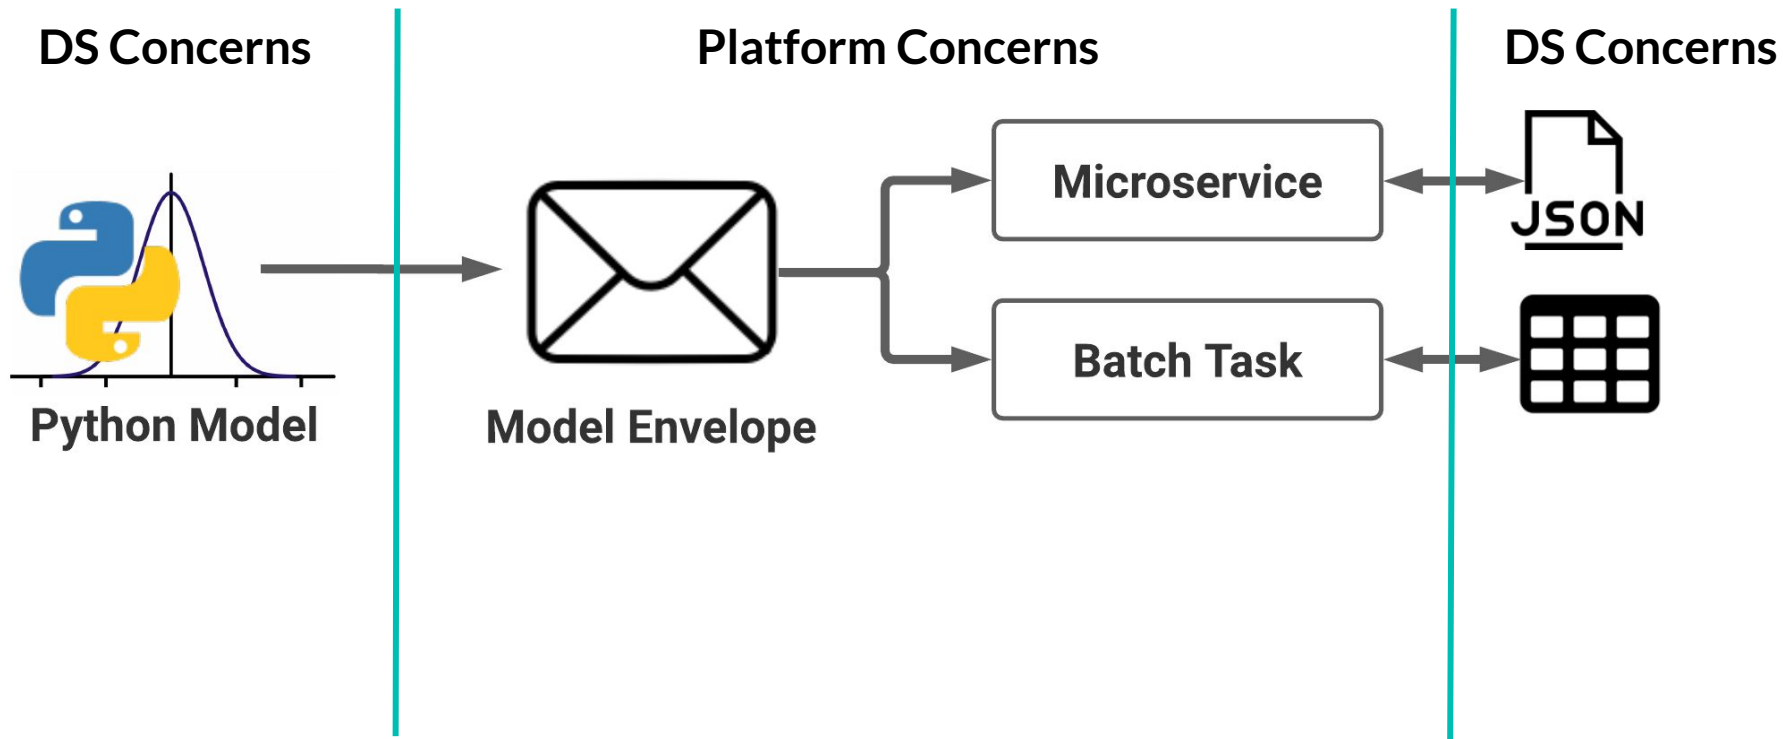
# Model Deployment at Stitch Fix



Once a model is in  
an envelope...

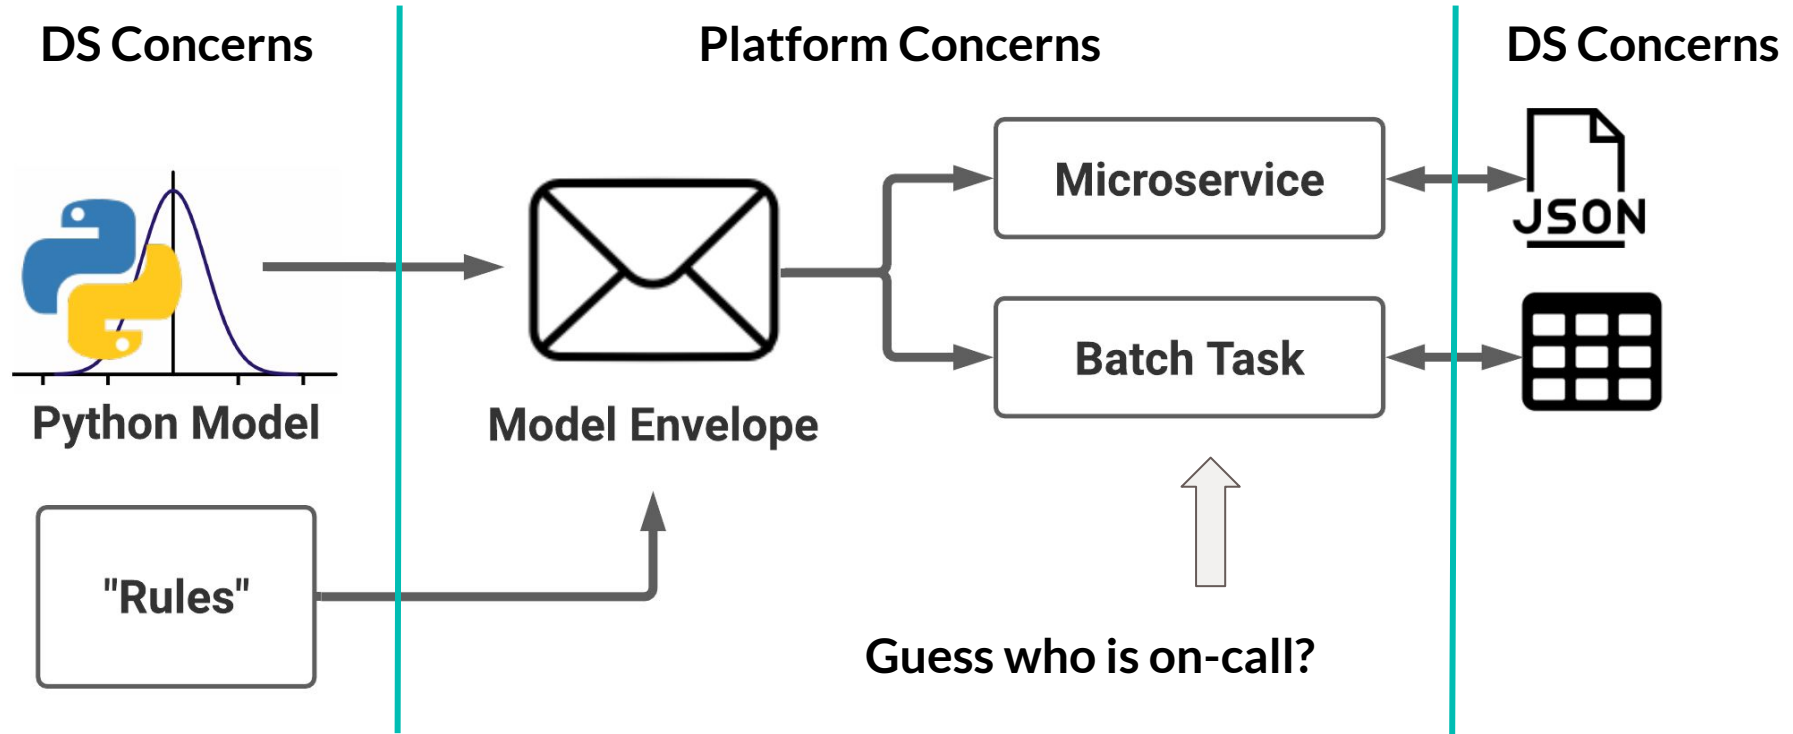
This comes for free!

# Who owns what?



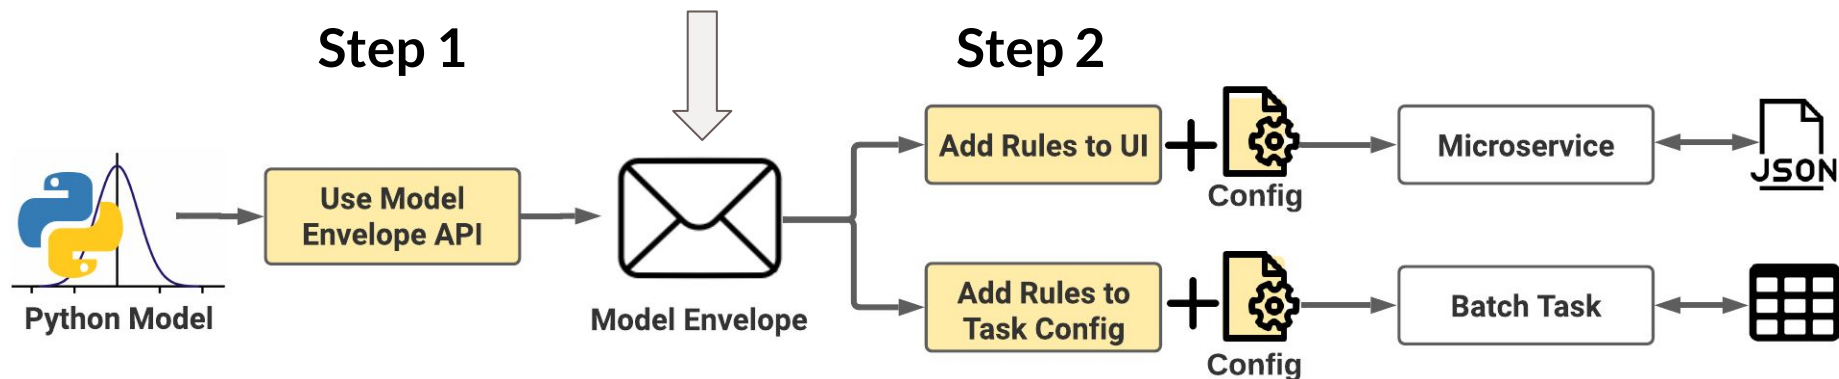


# Deployments are “triggered”



# Reality: two steps to get a model to production

Can be a terminal point.



**Self-service: takes <1 hour**  
**No code is written!**

# Step 1. save a model via Model Envelope API

etl.py

```
import model_envelope as me
from sklearn import linear_model
```

```
df_X, df_y = load_data_somewhat()
model = linear_model.LogisticRegression(multi_class='auto')
model.fit(df_X, df_y)
```

```
my_envelope = me.save_model(instance_name='my_model_instance_name',
                             instance_description='my_model_instance_description',
                             model=model,
                             query_function='predict',
                             api_input=df_X, api_output=df_y,
                             tags={'canonical_name': 'foo-bar'})
```

Note: no deployment trigger in ETL code.

# Step 2a. deploy model as a microservice

## Go to Model Envelope Registry UI:

- 1) Create deployment configuration.
- 2) Create **Rule** for auto deployment.
  - a) Else query for model & hit deploy.
- 3) Done.

Result:

- Web service with API endpoints
  - Comes with a Swagger UI & schema
- Model in production < 1 hour.

The screenshot displays the Model Envelope Registry UI. At the top, there's a search bar with 'skrawczyk' entered and a 'GO!' button. Below this, a table lists models, with one model 'abc\_e2e\_kmeans\_as\_of\_20210423' highlighted. To the right of the table, there are buttons for 'skrawczyk's models' and 'model\_lifecycle's models'. Below the table, a section titled 'API Endpoints' lists four endpoints: 'GET / Root', 'POST /v0/query V0 Query', 'POST /v0/query\_bulk V0 Query Bulk', and 'GET /v0/model\_info V0 Model Info'.

Diff	Model
<input type="checkbox"/>	<a href="#">abc_e2e_kmeans_as_of_20210423</a>

Method	Endpoint	Description
GET	/ Root	
POST	/v0/query	V0 Query
POST	/v0/query_bulk	V0 Query Bulk
GET	/v0/model_info	V0 Model Info


## Step 2b. deploy model as a batch task

### Create workflow configuration:

- 1) Create batch inference task in workflow.
  - a) Specify **Rule** & inputs + outputs.
- 2) Deploy workflow.
- 3) Done.

Result:

- Spark or Python task that creates a table.
- We keep an inference log.
- Model in production < 1 hour.



A large, faint, circular watermark of the Stitch Fix logo is visible on the left side of the slide. It features a stylized 'S' and 'F' inside a circle, with a grid pattern.

Stitch Fix

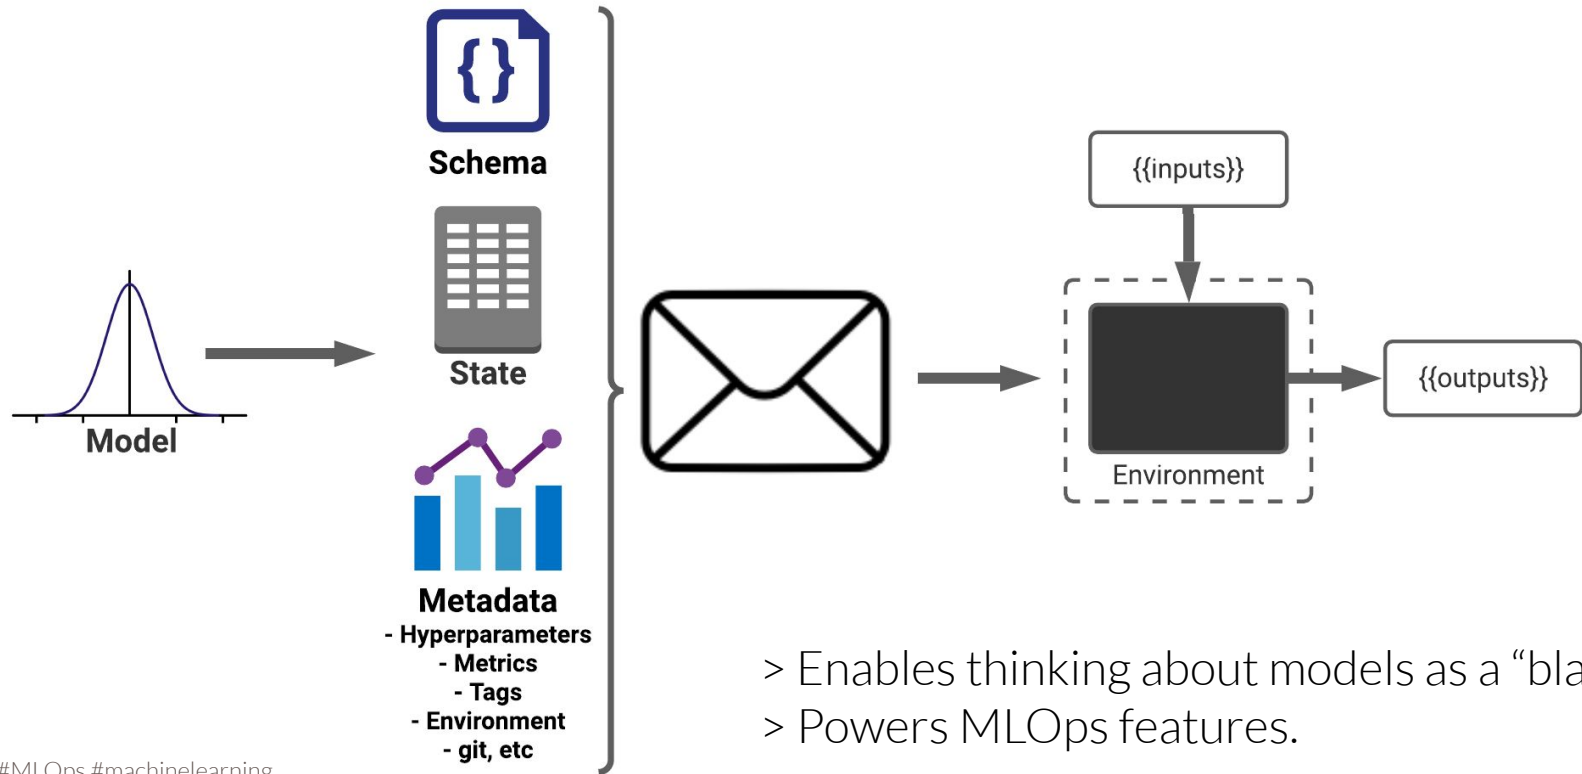
“Deployment for free”

> Model Envelope & envelope mechanics

Impact of being on-call

Summary & Future Work

# Q: What is the Model Envelope? A: It's a container.



# Wait this feels familiar?

**You:** “MLFlow/Verta much?”

**Me:** Yes & No.

This is all internal code -- nothing from open source.


In terms of functionality we're closer to a mix of:

- [MLFlow](#) | [Verta.ai](#)
- [ModelDB](#)
- [TFX](#)

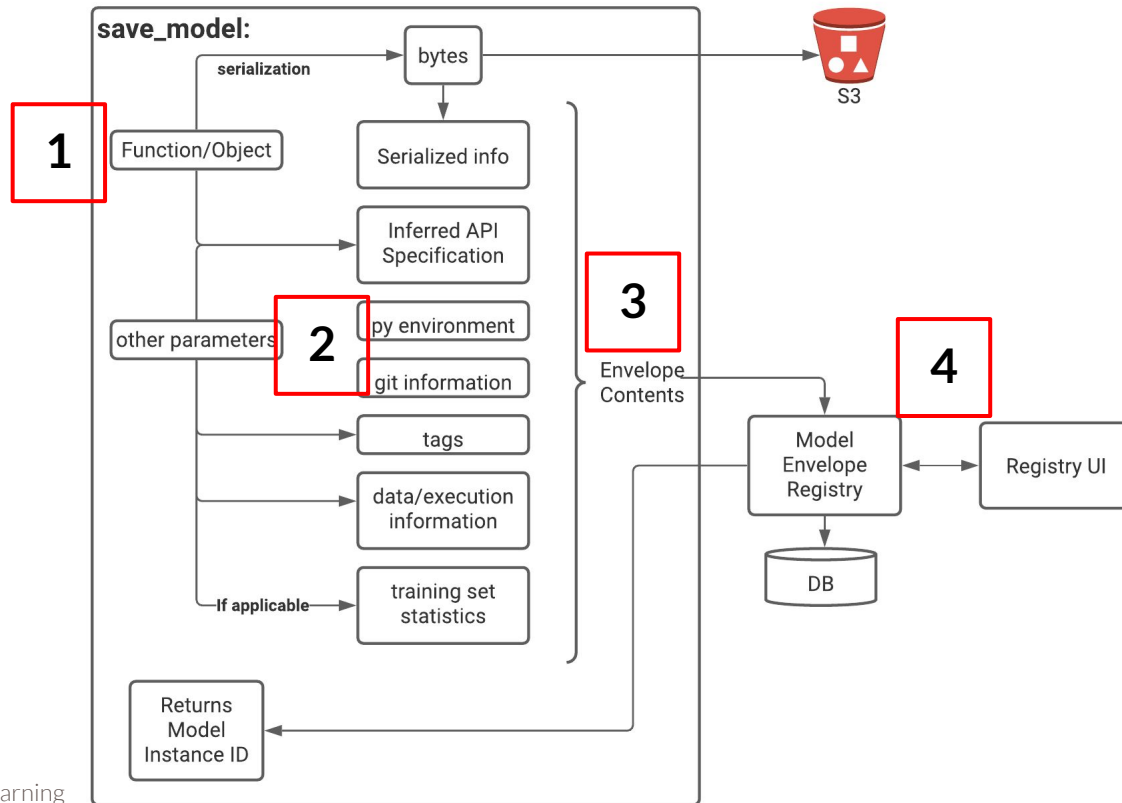
But this talk is too short to cover everything...



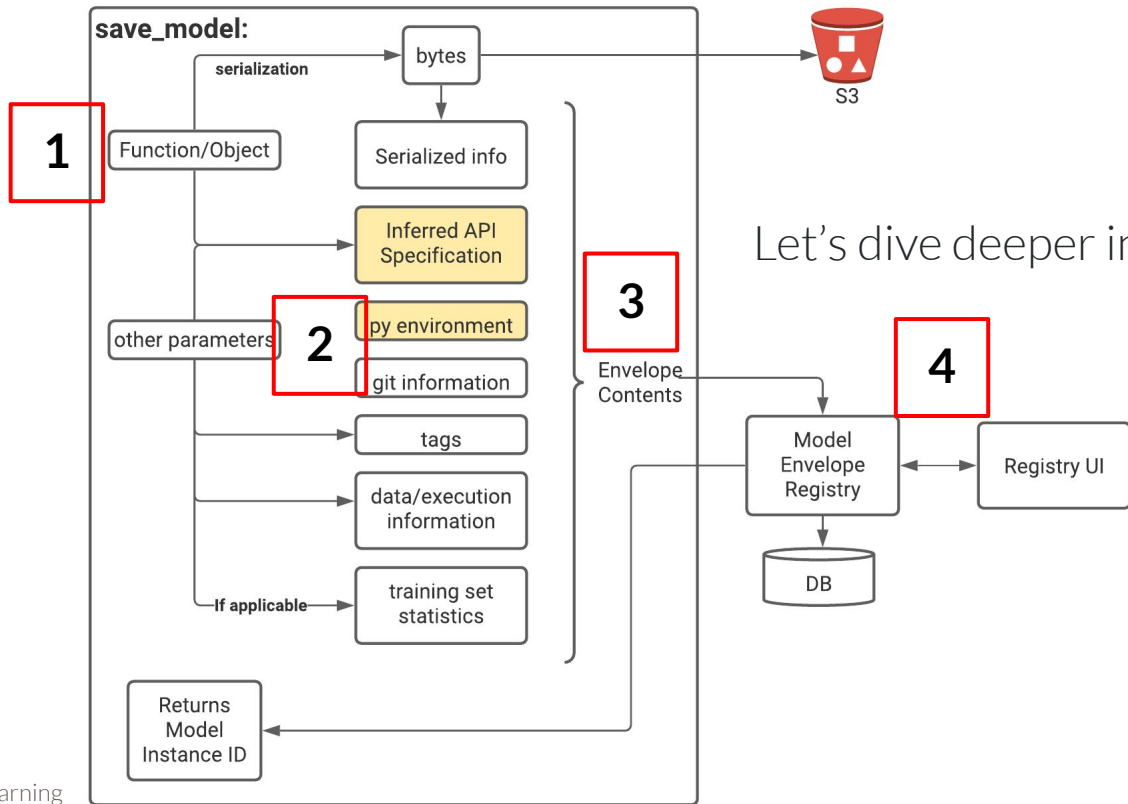
# Typical Model Envelope use

1. call **save\_model()** right after model creation in an ETL.
2. also have APIs to save metrics & hyperparameters, and retrieve envelopes.
3. once in an  **information is immutable** *except*:
  - a. tags -- for curative purposes.
  - b. metrics -- can add/adjust metrics.

# What does `save_model()` do?



# What does `save_model()` do?



# How do we infer a Model API Schema?

**Goal:** infer from code rather than explicit specification.

Require either fully annotated functions with only python/typing standard types:

```
def good_predict_function(self, x: float, y: List[int]) -> List[float]:  
def predict_needs_examples_function(self, x: pd.DataFrame, y):
```

Or, example inputs that are inspected to get a schema from:

```
my_envelope = me.save_model(instance_name='my_model_instance_name',  
                             instance_description='my_model_instance_description',  
                             model=model,  
                             query_function='predict',  
                             api_input=df_x, api_output=df_y,  
                             tags={'canonical_name': 'foo-bar'})
```

required for DF inputs →

# How do we infer a Model API Schema?

**Goal:** infer from code rather than explicit specification.

Require either fully annotated functions with only python/typing standard types:

## Why get a schema?

> Required for any form of validation:  
E.g. did the model get passed the right inputs?

## Why this way?

> To avoid breakage when something is updated.

```
def good_pr  
def predict
```

Or, example

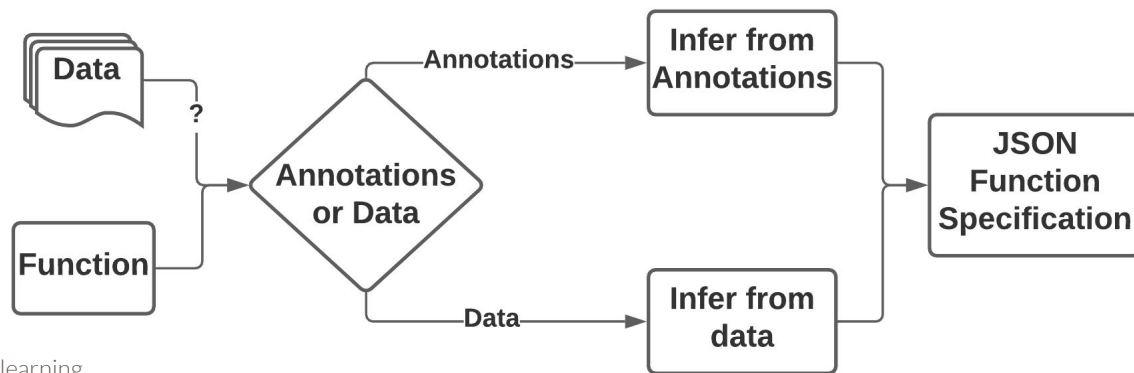
```
my_envelope
```

required for DF inputs →

```
model=model,  
query function='predict',  
api_input=df_x, api_output=df_y,  
tags={'canonical_name': 'foo-bar'})
```

# Model API Schema - Under the hood

- One of the most complex parts of the code base (90%+ test coverage!)
- We make heavy use of the *typing\_inspect* module & *isinstance()*.
  - We create a schema similar to TFX.
- Key component to enable exercising models in different contexts.
  - Enables code creation and input/output validation.
- Current limitations: **no default values in functions.**



# How do we capture python dependencies?

```
import model_envelope as me
from sklearn import linear_model

df_X, df_y = load_data_somewhat()
model = linear_model.LogisticRegression(multi_class='auto')
model.fit(df_X, df_y)

my_envelope = me.save_model(instance_name='my_model_instance_name',
                             instance_description='my_model_instance_description',
                             model=model,
                             query_function='predict',
                             api_input=df_X, api_output=df_y,
                             tags={'canonical_name': 'foo-bar'})
```

Point: no explicit passing of scikit-learn to save\_model().

# How do we capture python dependencies?

```
import model_envelope as me  
from sklearn import linear_model
```

## Why auto capture dependencies?

- > Want to be able to reproduce & reuse models.
- > Easy for the user to get wrong.

```
model=model,  
query_function='predict',  
api_input=df_X, api_output=df_y,  
tags={'canonical_name':'foo-bar'})
```

Point: no explicit passing of scikit-learn to save\_model().



# How do we capture python dependencies?

## **Assumption:**

We all run on the same\* base linux environment in training & production.

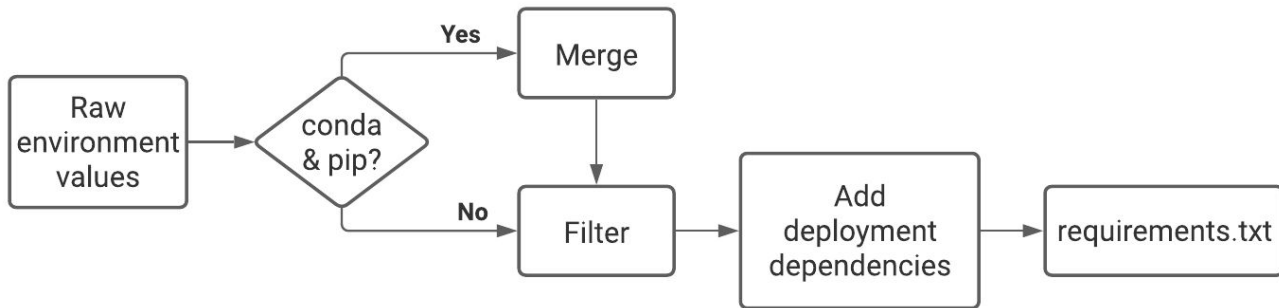
## **Store the following in the Model Envelope:**

- Result of `import sys; sys.version_info`
- Results of `> pip freeze`
- Results of `> conda list --export`

## **Local python modules** (not installable):


- Add modules as part of `save_model()` call.
- We store them with the model bytes.

# How do we build the python deployment env.?



## Filter:

- hard coded list of dependencies to filter. E.g. jupyterhub.
- upkeep cheap; add/update every few months.

A large, faint, circular watermark of the Stitch Fix logo is visible in the background on the left side of the slide. The logo consists of a circle with a grid of lines inside, forming a stylized 'S' or a basket-like shape.

# Stitch Fix

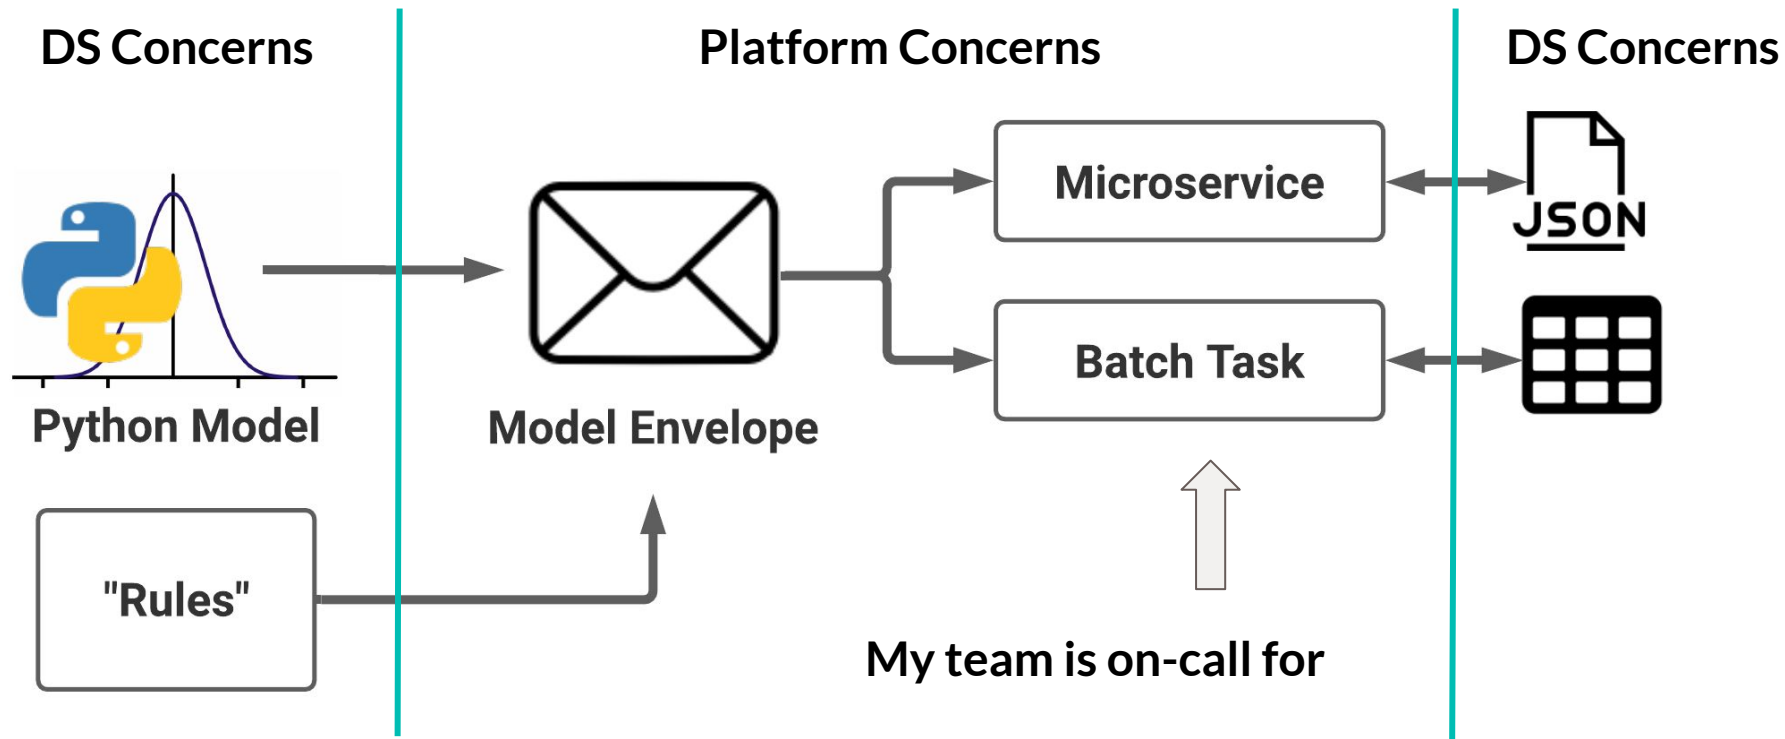
## “Deployment for free”

### Model Envelope & envelope mechanics

- > Impact of being on-call

## Summary & Future Work

# Remember this split:



# Impact of being on-call

## Two truths:

- No one wants to be paged.
- No one wants to be paged for a model they didn't write!



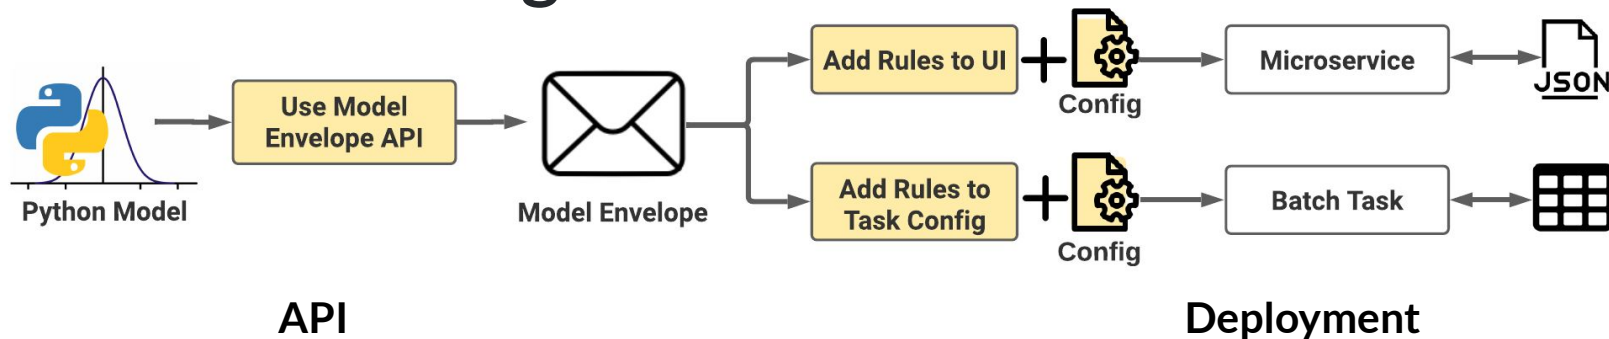
a pager

## But, this incentivizes Platform to build out MLOps capabilities:

- Capture bad models before they're deployed!
- Enable observability, monitoring, and alerting to speed up debugging.

Luckily we have autonomy and freedom to do so!

# What can we change?



Automatic capture == license to change:

- Model API schema
- Dependency capture
- Environment info: git, job, etc.

Incentives for DS to additionally provide:

- Datasets for analysis
- Metrics
- Tags

MLOps approaches to:

- Model validation
- Model deployment & rollback
- Model deployment vehicle:
  - From logging, monitoring, alerting
  - To architecture: microservice, or Ray, or?
- Dashboarding/UIs

# Overarching benefit

1. Data Scientists get to focus more on modeling.
  - a. more business wins.
2. Platform focuses on MLOps:
  - a. can be a rising tide that raises all boats!

A large, faint, circular watermark logo for Stitch Fix is visible in the background on the left side of the slide. It features a stylized 'S' and 'F' intertwined within a circle.

Stitch Fix

“Deployment for free”

Model Envelope & envelope mechanics

Impact of being on-call

> Summary & Future Work



# Summary - “Deployment for free”

## We enable deployment for free by:

- Capturing a comprehensive model artifact we call the Model Envelope.
- The Model Envelope facilitates code & environment generation for model deployment.
- Platform owns the Model Envelope and is on-call for generated services & tasks.

## Business wins:



- Data Scientists get to focus more on modeling.
- Platform is incentivized to improve and iterate on MLOps practices.

# Future Work

- **Better MLOps features:**
  - Observability, scalable data capture (e.g. whylogs), & alerting.
  - Model Validation & CD patterns.
- **“Models on Rails”:**
  - Target specific SLA requirements.
- **Configuration driven model creation:**
  - Abstract away glue code required to train & save models.



# Thank you! We're hiring! Questions?

 @stefkrawczyk  
 [linkedin.com/in/skrawczyk](https://www.linkedin.com/in/skrawczyk)

Try out Stitch Fix → [goo.gl/Q3tCQ3](https://goo.gl/Q3tCQ3)

#CS329S #MLOps #machinelearning

STITCH FIX